

Report

NS2 Term Project: *TCP Vegas-A: Improving the Performance of TCP Vegas*



Date of Submission: 28th February 2023

Submitted by:

Sanjida Islam Era
1805116
L-3/T-2
CSE-B2

Introduction:

Congestion control is a crucial aspect of network performance, especially in wired and wireless networks, where resources are limited, and network traffic is high. Among the several new features implemented in TCP Vegas that was originally proposed by Brakmo et al., one of the most important differences between TCP Vegas and TCP Reno is its congestion avoidance scheme. While TCP Reno (and its variants like NewReno) rely on packet loss detection to detect network congestions, TCP Vegas uses a sophisticated bandwidth estimation scheme to proactively gauge network congestion. While it has been shown that TCP Vegas provides better performance compared to TCP Reno, studies have identified various issues associated with the protocol. In this project, we propose modifications to the congestion avoidance mechanism of the TCP Vegas to overcome these limitations which is proposed in the paper " TCP Vegas-A: Improving the Performance of TCP Vegas" by K.N. Srijith, Lillykutty Jacob, A.L. Ananda and a modification of this.

The proposed congestion control algorithm aims to improve TCP throughput and some other parameters like delivery ratio in wired and wireless networks by obtaining a fairer share of the network bandwidth in wired and satellite scenarios. The algorithm achieves this by dynamically adjusting the congestion window size of each TCP Vegas based on network conditions, and probing the network for available bandwidth. We have also modified this Vegas-A algorithm to achieve higher performance in wired networks.

To evaluate the performance change, we have compared the metrics of the modified algorithm to the existing algorithm of TCP Vegas, and analysed the corresponding graphs subsequently.

Network Topologies:


The network topology assigned to this project were:

1. Wired nodes
2. Wireless mobile nodes with MAC 802.11

Here the assigned topology was "Wireless mobile nodes with MAC 802.15.4" but it results results negative throughput. Since no possible solution to this problem was found at the time of concluding data collection, the agent type was switched to IEEE 802.11.

We have performed necessary simulations and derived results accordingly of these two kinds of networks. Visualisations of both the networks are as follows:

Wired nodes:



Wireless mobiles nodes:



Parameters under variation:

The parameters varied for the topologies in the simulation are as follows:

1. Number of nodes: 20, 40, 60, 80, 100
2. Number of flows: 10, 20, 30, 40, 50
3. Number of packets sent per second: 100, 200, 300, 400, 500
4. Speed of mobile nodes: 5 m/s, 10 m/s, 15 m/s, 20 m/s, 25 m/s
(applicable for the simulation involving mobile nodes in wireless network)

Modifications made:

The modification of this project is regarding the Congestion Avoidance mechanism of TCP Vegas. The reference paper presents modification to TCP Vegas and refers to the modified algorithm as TCP Vegas-A, where 'A' stands for adaptive. TCP Vegas uses fixed values for variables a and b , set to 1 and 3 usually. Recall that Vegas's strategy is to adjust the source's congestion window in an attempt to keep a small number of packets buffered in the routers along the path. Still subscribing to the idea that the average number of packets in the router buffer is to be kept within a and b , the main idea behind TCP Vegas-A is that rather than fixing a and b , they be made dynamically changeable, adaptive. At the start of a connection, a is set to 1 and b to 3. These values are then changed dynamically depending on the network conditions. Another way of looking at this modification is that we are trying to bring the network probing capability of TCP Reno into TCP Vegas

The congestion control in TCP Agents are divided into three phases:

1. Slow start: The initial phase of TCP connection establishment in which the sender starts transmitting data to the receiver with a conservative transmission rate. During this phase, the sender increases its transmission rate exponentially until it detects congestion in the network.
2. Congestion Avoidance: The slow start threshold ($ssthresh$) is reached, and we increase the congestion window according to the protocol.
3. Congestion Control: Congestion is detected, and similar to avoidance, a procedure is followed to decrease the window to decrease the loss due to congestion.

These phases are handled by two functions in ns2, namely `openwnd()` and `slowdown()`. While slow start and congestion recovery algorithms of Vegas-A are the same as that of Vegas, we propose a modified congestion avoidance mechanism.

I implemented the following changes:

In tcp.h:

1. New variables for optimal congestion window calculation:
Specially for calculating the

```

tcp-vegas.cc 4  C tcp.h 3 x
C tcp.h > ...
611
612 virtual void delay_bind_init_all();
613 virtual int delay_bind_dispatch(const char *varName, const char *localName, TclObject *tracer);
614
615 double t_cwnd_changed_; // last time cwnd changed
616 double firstrecv_; // time recv the 1st ack
617
618 int v_alpha_; // vegas thruput thresholds in pkts
619 int v_beta_;
620
621 int v_gamma_; // threshold to change from slow-start to
622 // congestion avoidance, in pkts
623
624 int v_slowstart_; // # of pkts to send after slow-start, deflt(2)
625 int v_worried_; // # of pkts to chk after dup ack (1 or 2)
626
627 double v_timeout_; // based on fine-grained timer.
628 double v_rtt_;
629 double v_sa_;
630 double v_sd_;
631 double tht_prev=0; // newly added
632
633 int v_cntRTT_; // # of rtt measured within one rtt
634 double v_sumRTT_; // sum of rtt measured within one rtt
635
636 double v_begtime_; // tagged pkt sent
637 int v_begseq_; // tagged pkt seqno
638
639 double* v_sendtime_; // each unacked pkt's sendtime is recorded.
640 int* v_transmits_; // # of retx for an unacked pkt
641
642 int v_maxwnd_; // maxwnd size for v_sendtime[]

```

In tcp-vegas.cc:

1. Modified functions to run when modifications are not enabled:

Mainly I have done this modification in `recv()` function of `tcp-vegas.cc`. I have added an extra part in the congestion avoidance section. For this I have declared some parameters to calculate current throughput using current rtt, and previous throughput by the previous declared parameter `tht_prev`. Other local variables are given:

```

// expect = (current window size)/baseRTT
expect = double(t_seqno_-last_ack_)/v_baseRTT_;
double new_actual = double(t_seqno_-last_ack_)/rtt;

// calc actual and expect thruput diff, delta
int delta=int((expect-v_actual_)*v_baseRTT_+0.5);
double delta_check = (expect-new_actual);
int cwnd_prev = cwnd_;
int ssthresh_prev = ssthresh_;

```

2. The main modification for Vegas-A:

```

tcp-vegas.cc 3 x  tcp.h 3
tcp-vegas.cc > recv(Packet *, Handler *)
817 // printf("actual - %f\n", new_actual);
818 // printf("expected: %f\n", expect);
819
820 if(cwnd_prev >= ssthresh_prev){
821     if(delta_check < v_beta_ && delta_check > v_alpha_){
822         if(new_actual>tht_prev){
823             v_alpha=v_alpha+1;
824             v_beta=v_beta+1;
825         }
826     }
827     else if(delta_check<v_alpha_){
828         if(v_alpha>1){
829             if(new_actual < tht_prev){
830                 v_alpha=v_alpha-1;
831                 v_beta=v_beta-1;
832             }
833         }
834     }
835     else if(delta_check>v_beta_){
836         if(v_alpha>1){
837             v_alpha=v_alpha-1;
838             v_beta=v_beta-1;
839         }
840     }
841     else{
842
843     }
844 }
845 } // end of if(rtt > 0)

```

3. Another approach in modification for getting better result in wired topology (by modifying the Vegas-A algorithm):

```

tcp-vegas.cc 3 x  tcp.h 3
tcp-vegas.cc > recv(Packet *, Handler *)
817 // printf("actual - %f\n", new_actual);
818 // printf("expected: %f\n", expect);
819
820 if(cwnd_prev >= ssthresh_prev){
821     if(delta_check < v_beta_ && delta_check > v_alpha_){
822         if(new_actual>tht_prev){
823             v_alpha=v_alpha+1;
824             v_beta=v_beta+1;
825         }
826     }
827     else if(delta_check<v_alpha_){
828         if(v_alpha>1){
829             if(new_actual < tht_prev){
830                 v_alpha=v_alpha-1;
831             }
832         }
833     }
834     else if(delta_check>v_beta_){
835         if(v_alpha+1 < v_beta_){
836             v_beta=v_beta-1;
837         }
838     }
839     else{
840
841     }
842 }
843 } // end of if(rtt > 0)
844

```

4. Finally setting the throughput as previous throughput for next rtt:


```
861         if(rtt>0) tht_prev=v_actual_;  
862         else tht_prev=0;
```

Results with graphs:


1. Wireless 802.11 Mobile nodes:

With respect to Node Speed:


Network Throughput:




End-to-End Delay:




Packet Delivery Ratio:



Packet Drop Ratio:




Energy Consumption:




With respect to Number of Nodes:


Network Throughput:




End-to-End Delay:




Packet Delivery Ratio:




Packet Drop Ratio:




Energy Consumption:




With respect to Number of Flows:
Network Throughput:




End-to-End Delay:




Packet Delivery Ratio:



Packet Drop Ratio:




Energy Consumption:




With respect to Number of Packets per second:


Network Throughput:




End-to-End Delay:




Packet Delivery Ratio:



Packet Drop Ratio:




Energy Consumption:




2. Wired nodes:

With respect to Number of Nodes:


Network Throughput:



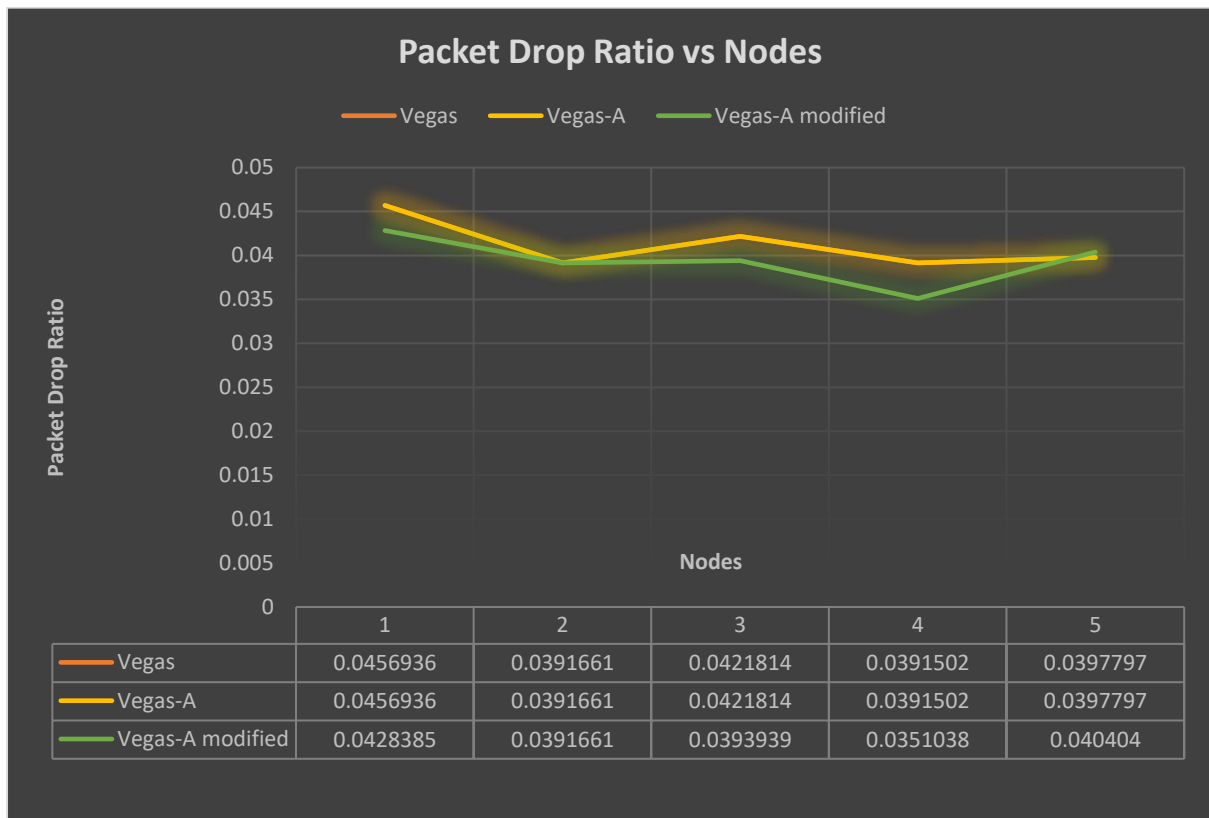
End-to-End Delay:




Packet Delivery Ratio:




Packet Drop Ratio:




With respect to Number of Flows:
Network Throughput:



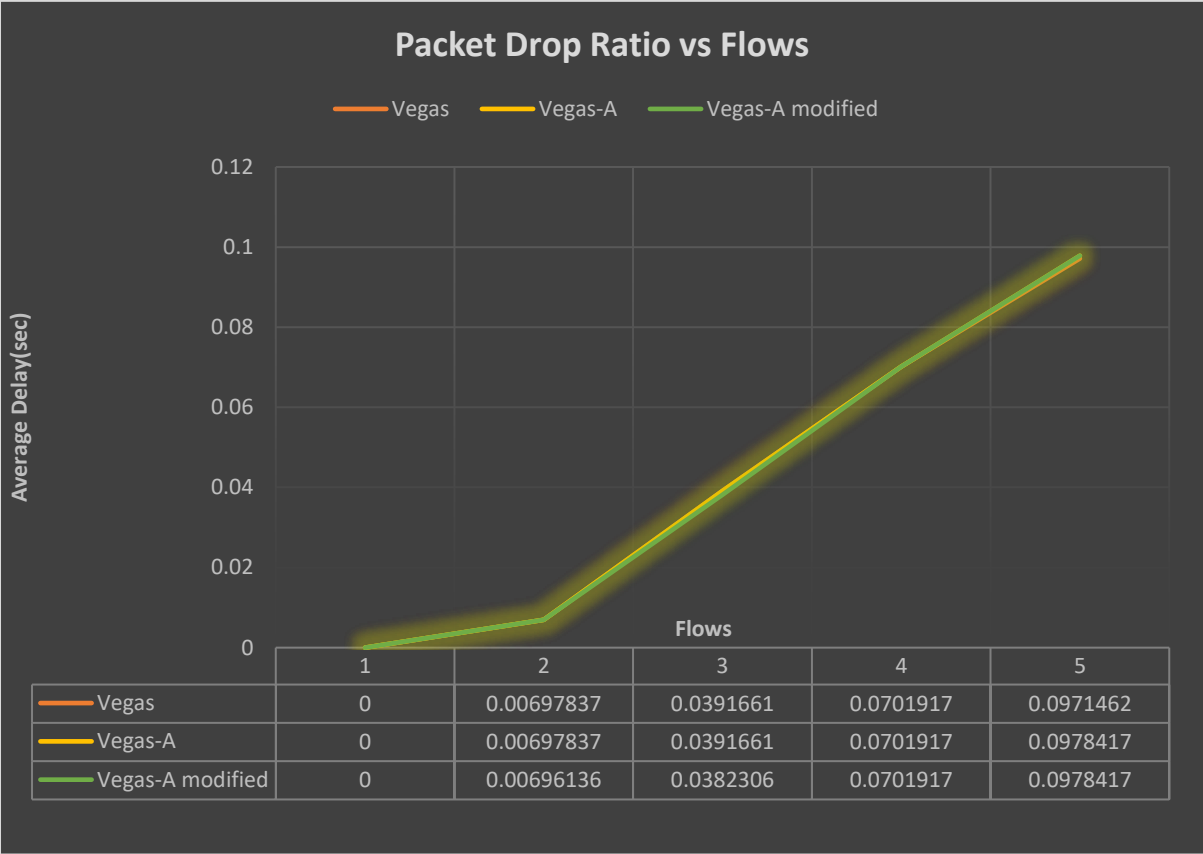
End-to-End Delay:




Packet Delivery Ratio:



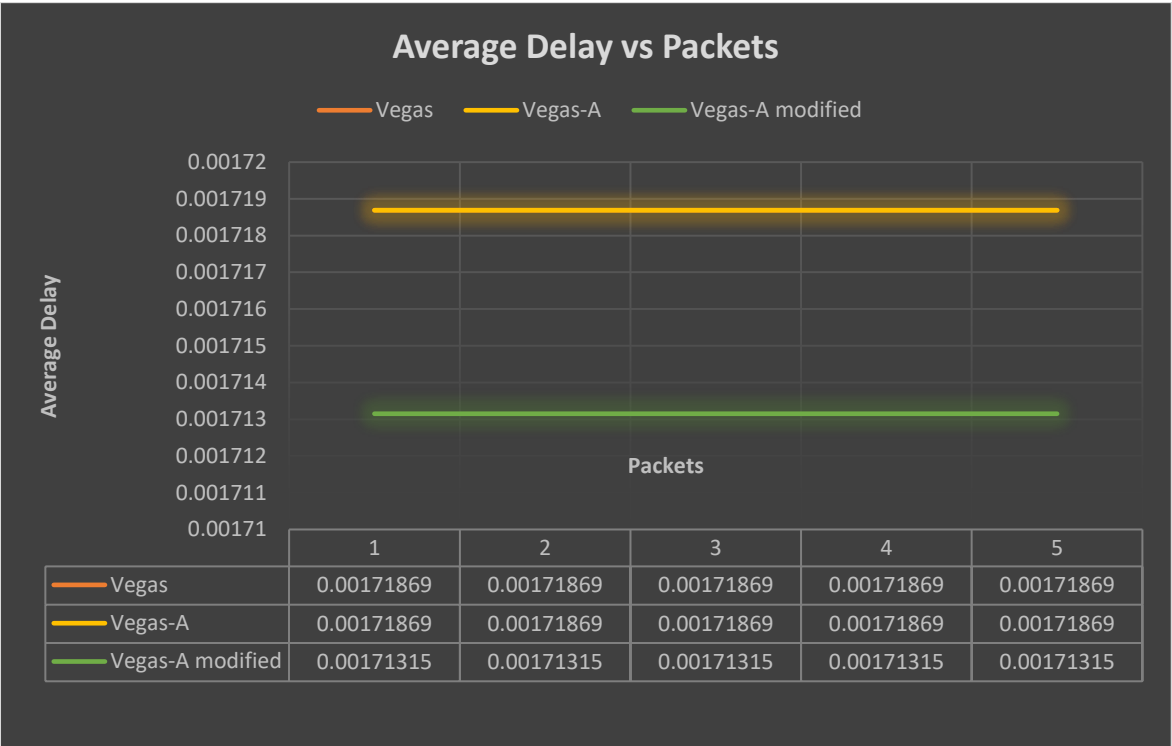
Packet Drop Ratio:




With respect to Number of Packets per second:
Network Throughput:




End-to-End Delay:



Packet Delivery Ratio:



Packet Drop Ratio:



Summary of Vegas-A vs Vegas:

TCP Vegas-A shows improvement over the tradition TCP Vegas in cases of:

- Throughput, packet delivery ratio and packet drop ratio for wireless
- Throughput (in some cases) in wired

TCP Vegas-A lacks behind in case of:

- Average Delay and Energy Consumption for wireless

Summary of Modified Vegas-A vs Vegas:

TCP Vegas-A(Modified) shows improvement over the tradition TCP Vegas in cases of:

- Throughput, Packet Delivery Ratio and Packet Drop Ratio, Energy Consumption while varying packets sent per second for wireless
- All four parameters almost (Throughput, Packet Delivery Ratio and Packet Drop Ratio, Average Delay) for wired

